

MapReduce

CS 351

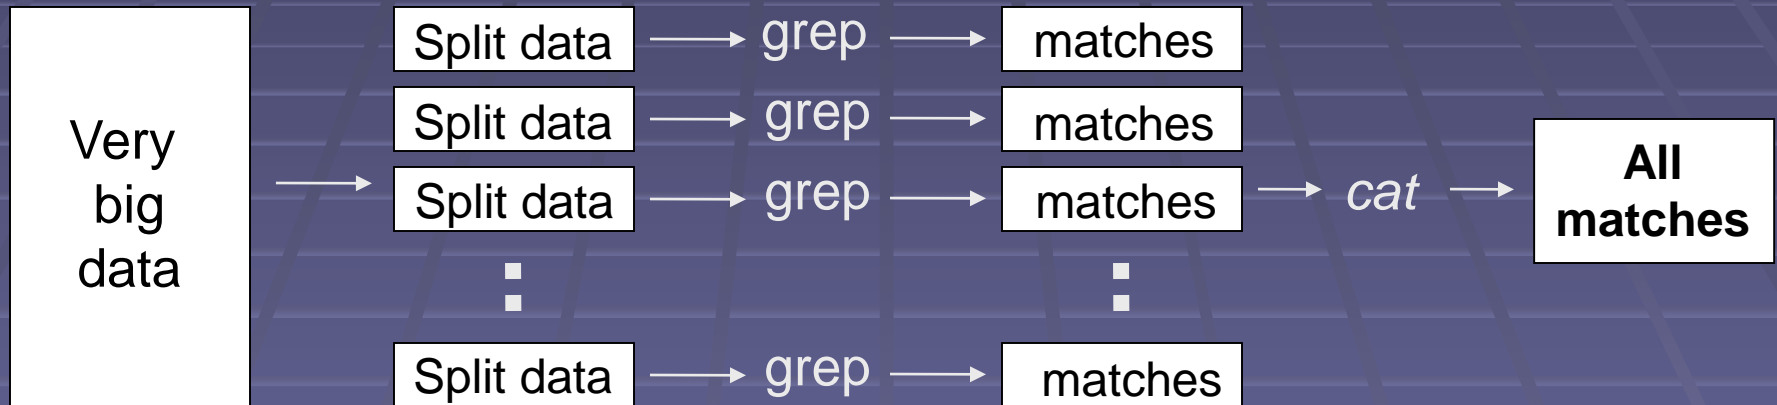
What is MapReduce

- Origin from Google, [OSDI'04]
- A simple programming model
- Functional model
- For large-scale data processing
 - Exploits large set of commodity computers
 - Executes process in distributed manner
 - Offers high availability

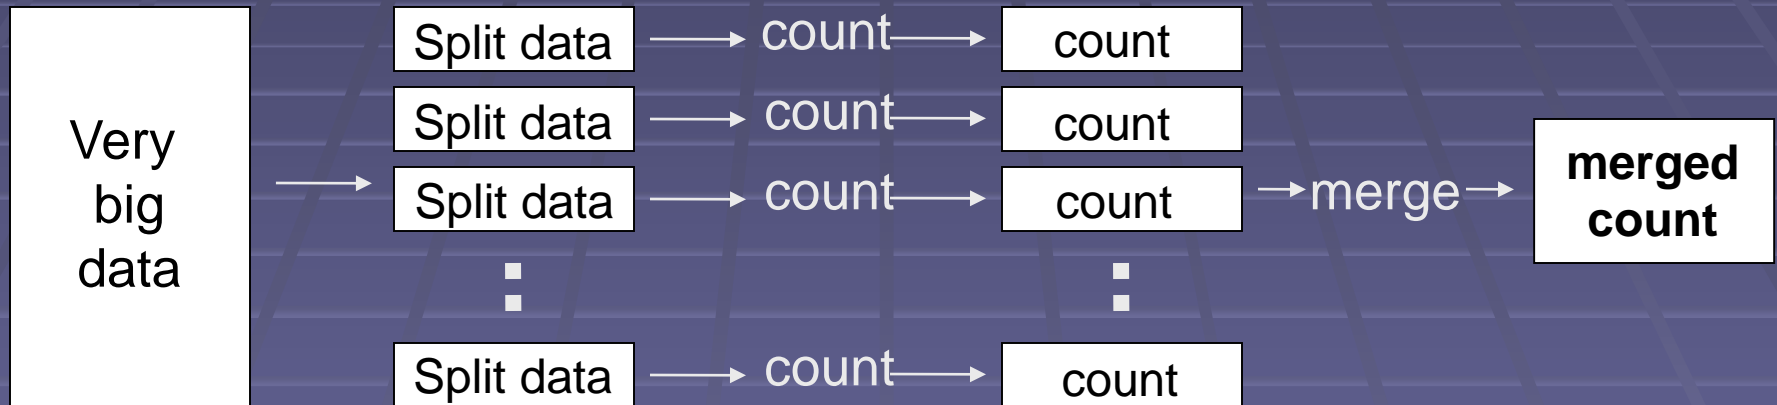
Motivation

- Lots of demands for very large scale data processing
- A certain common themes for these demands
 - Lots of machines needed (scaling)
 - Two basic operations on the input
 - Map
 - Reduce

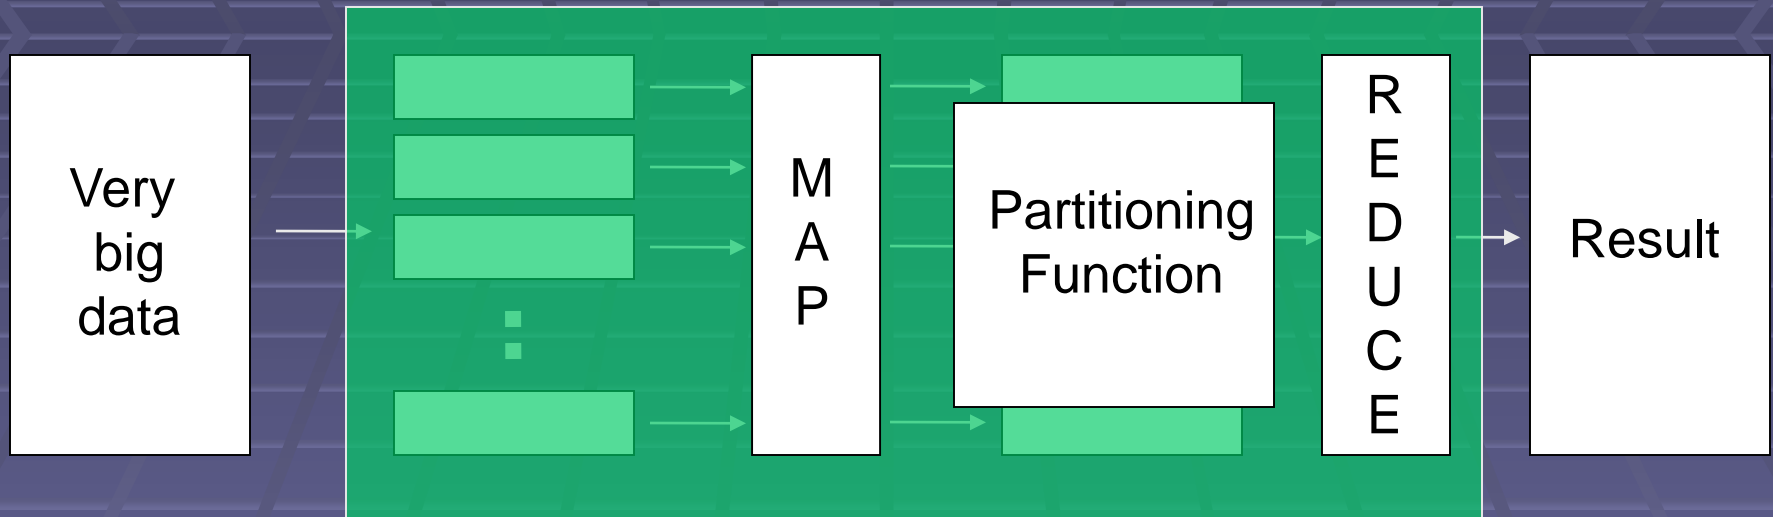
Distributed Grep



Distributed Word Count



Map+Reduce



- Map:

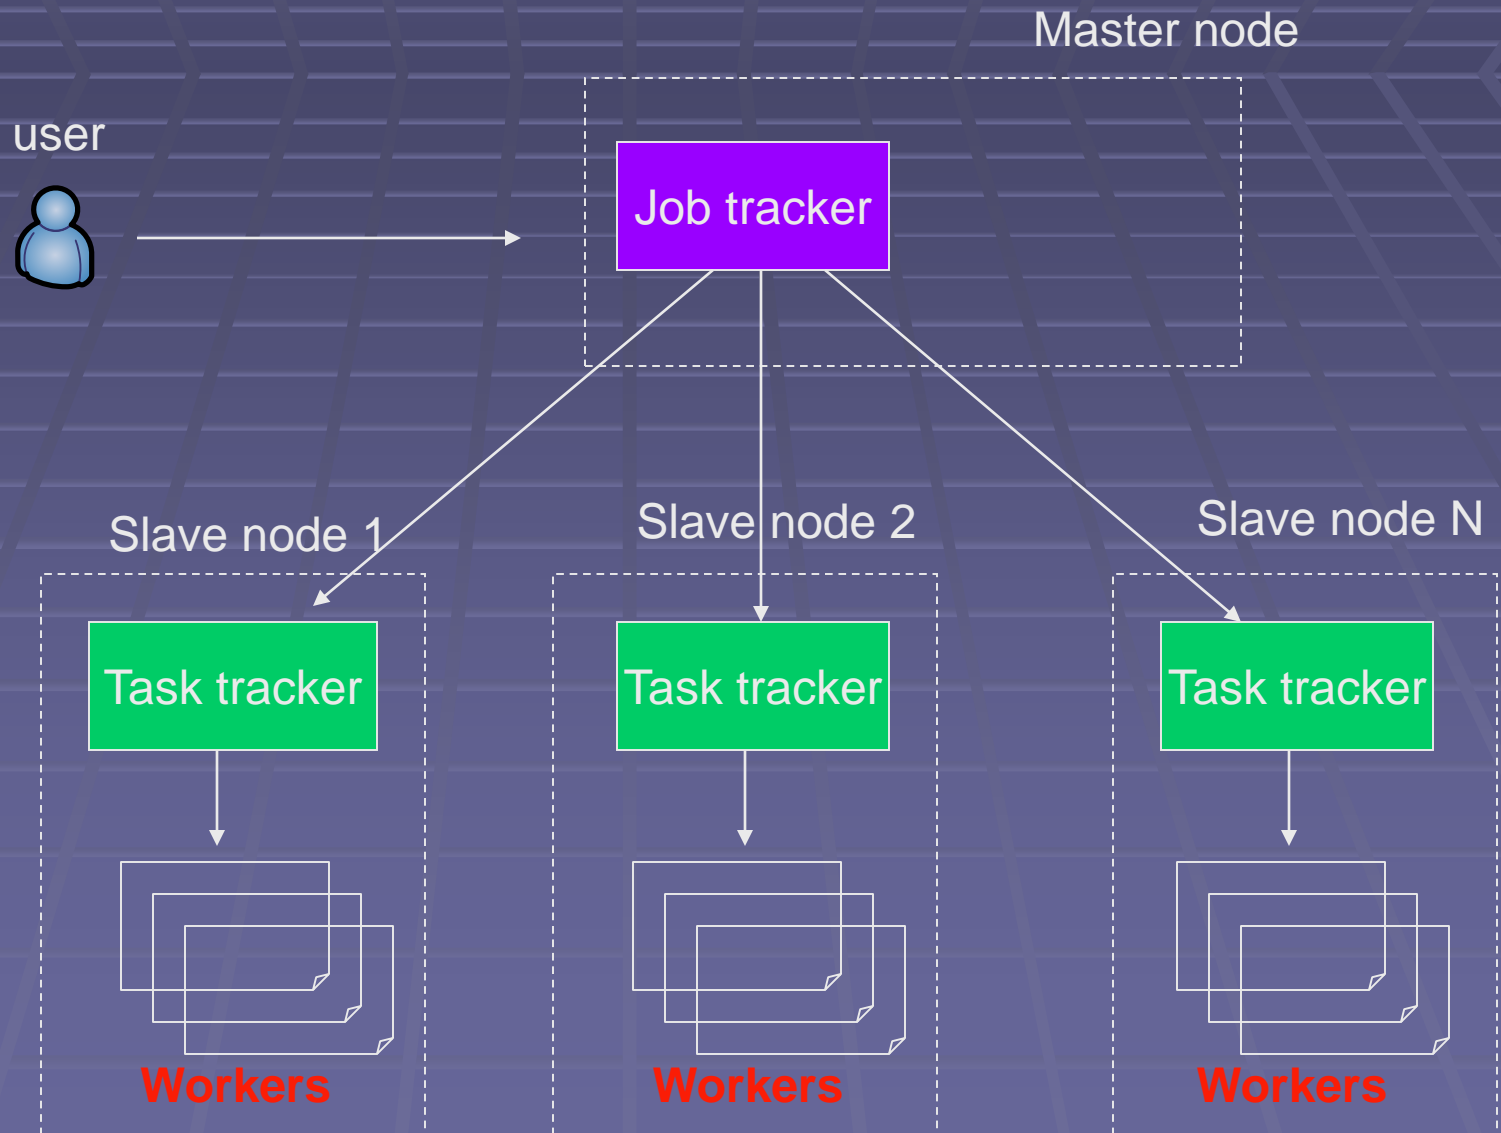
- Accepts *input* key/value pair
- Emits *intermediate* key/value pair

- Reduce :

- Accepts *intermediate* key/value* pair
- Emits *output* key/value pair

The design and how it works

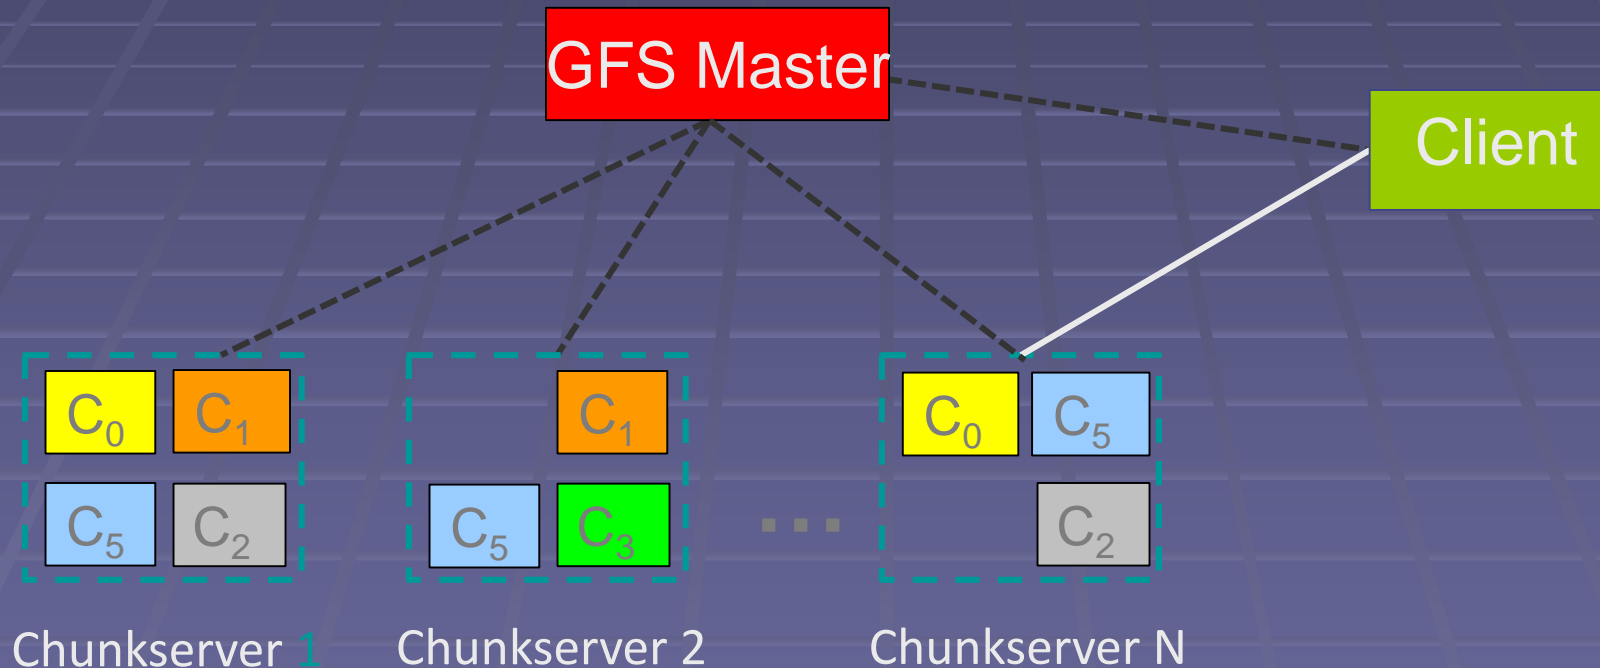
Architecture overview



GFS: underlying storage system

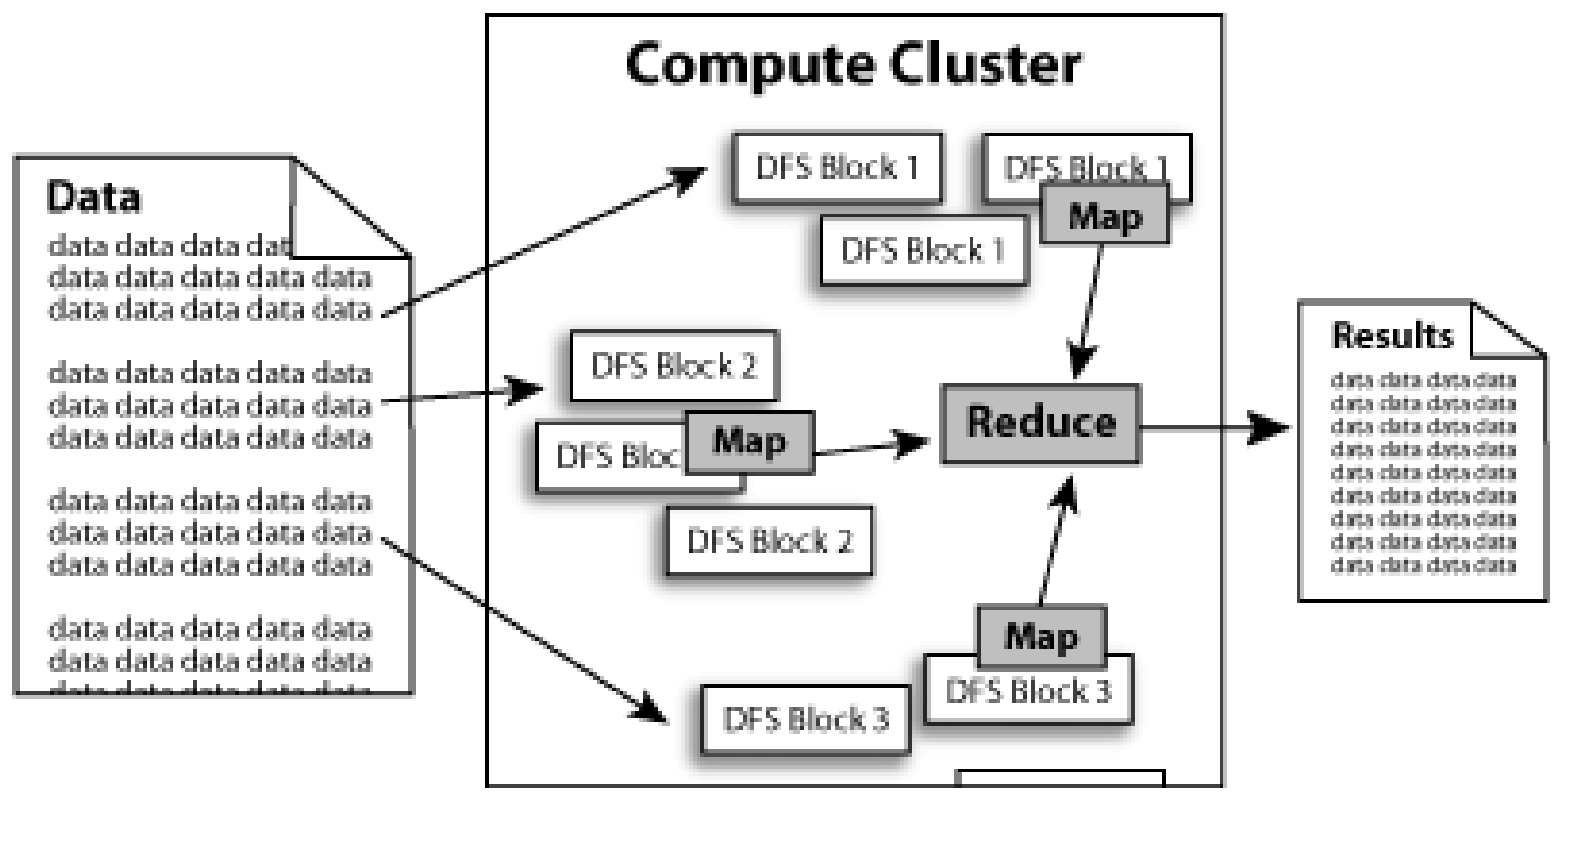
- Goal
 - global view
 - make huge files available in the face of node failures
- Master Node (meta server)
 - Centralized, index all chunks on data servers
- Chunk server (data server)
 - File is split into contiguous chunks, typically 16-64MB.
 - Each chunk replicated (usually 2x or 3x).
 - Try to keep replicas in different racks.

GFS architecture

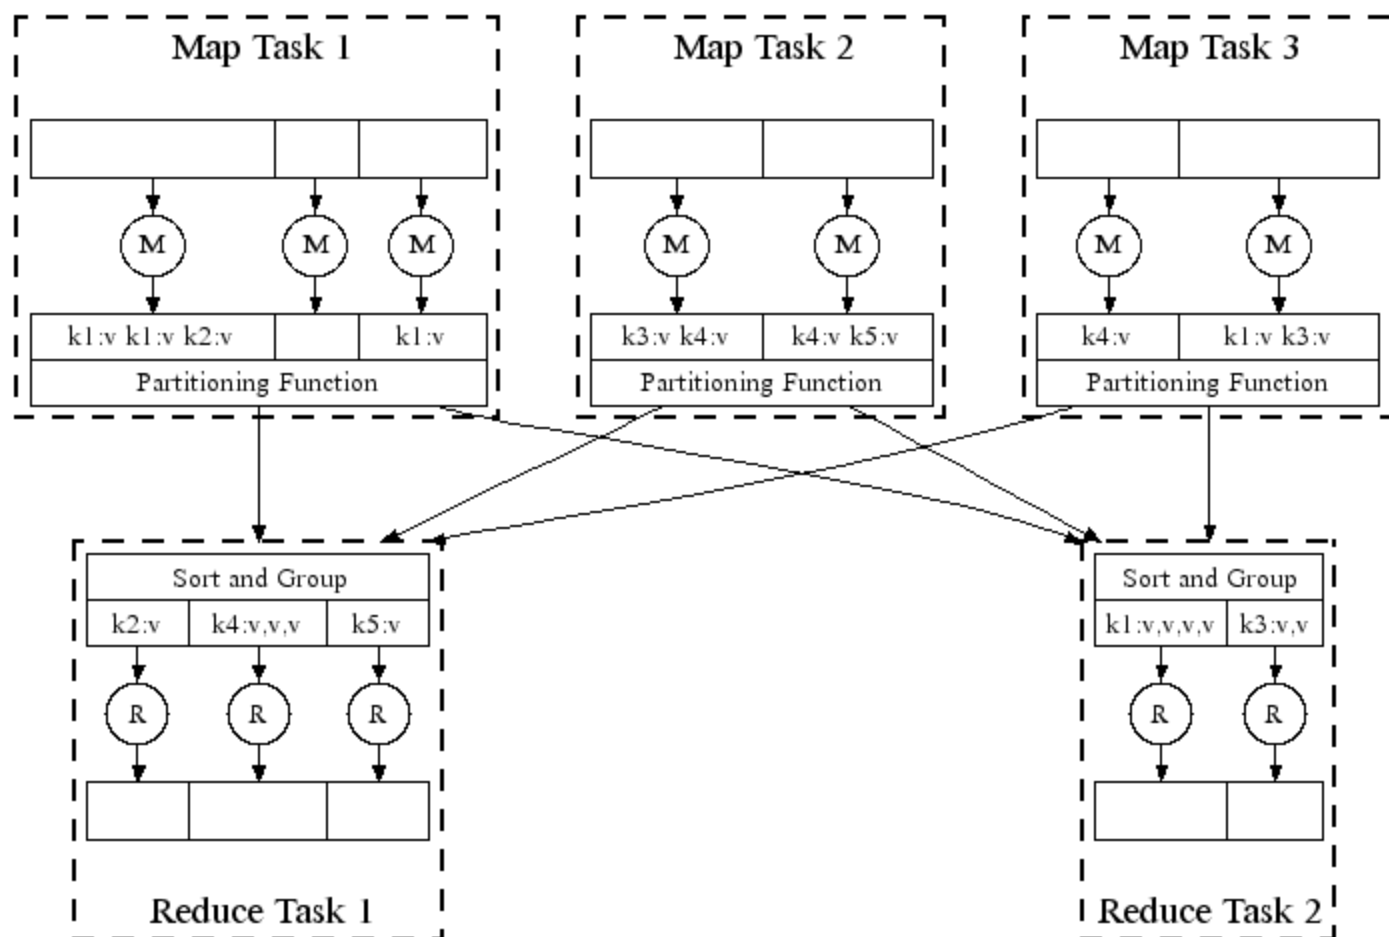


Functions in the Model

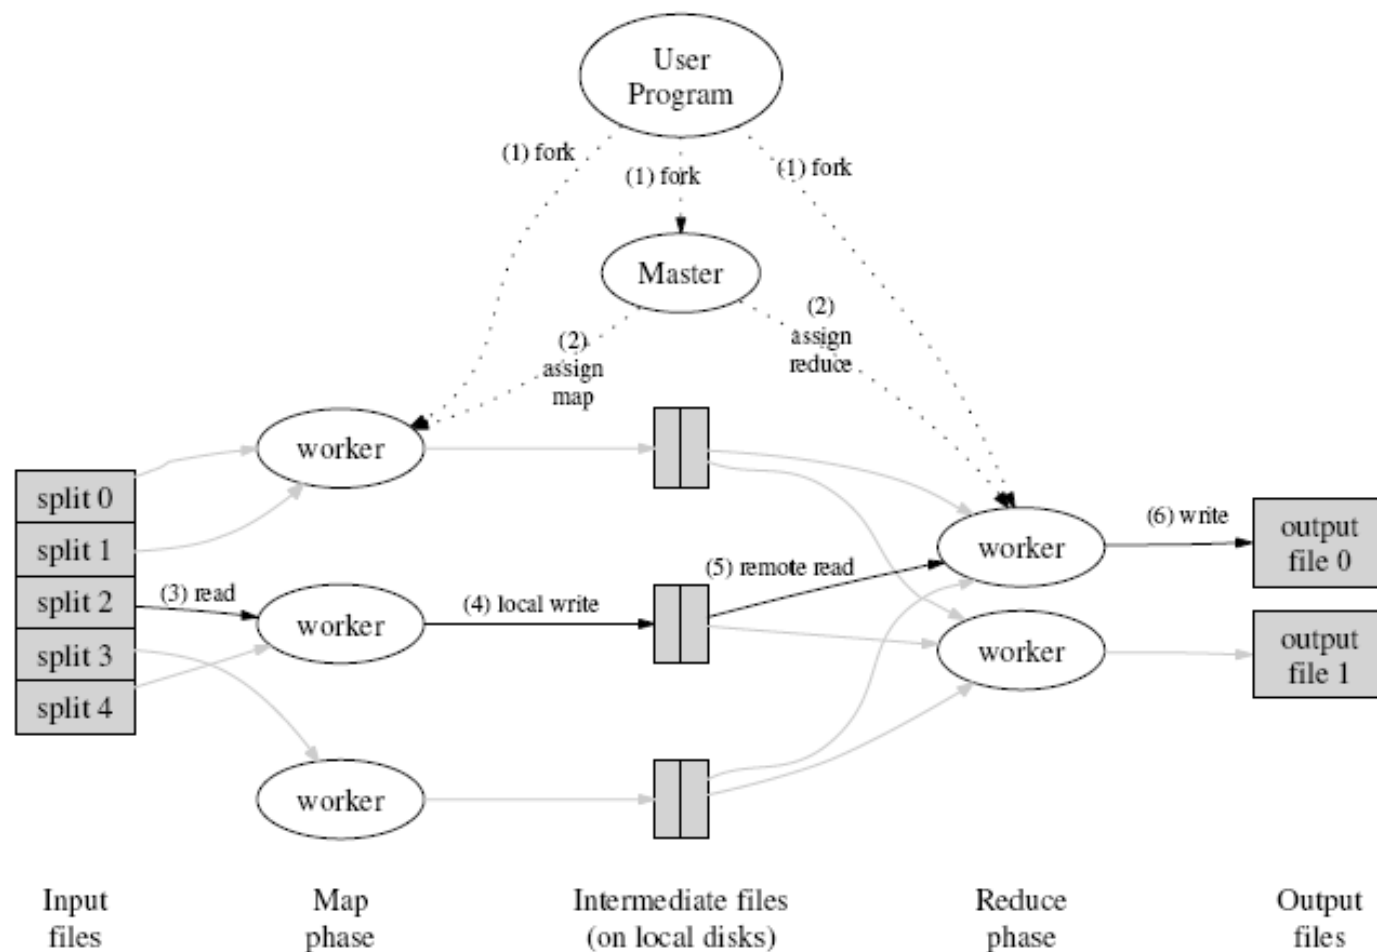
- Map
 - Process a key/value pair to generate intermediate key/value pairs
- Reduce
 - Merge all intermediate values associated with the same key
- Partition
 - By default : $\text{hash}(\text{key}) \bmod R$
 - Well balanced



Parallel Execution



How does it work?



Locality issue

- Master scheduling policy
 - Asks GFS for locations of replicas of input file blocks
 - Map tasks typically split into 64MB (== GFS block size)
 - Map tasks scheduled so GFS input block replica are on same machine or same rack
- Effect
 - Thousands of machines read input at local disk speed
 - Without this, rack switches limit read rate

Fault Tolerance

- Reactive way
 - Worker failure
 - Heartbeat, Workers are periodically pinged by master
 - NO response = failed worker
 - If the processor of a worker fails, the tasks of that worker are reassigned to another worker.
 - Master failure
 - Master writes periodic checkpoints
 - Another master can be started from the last checkpointed state
 - If eventually the master dies, the job will be aborted

Fault Tolerance

- Proactive way (**Redundant Execution**)
 - The problem of “stragglers” (slow workers)
 - Other jobs consuming resources on machine
 - Bad disks with soft errors transfer data very slowly
 - Weird things: processor caches disabled (!!)
 - When computation almost done, reschedule in-progress tasks
 - Whenever either the primary or the backup executions finishes, mark it as completed

Fault Tolerance

- Input error: bad records
 - Map/Reduce functions sometimes fail for particular inputs
 - Best solution is to debug & fix, but not always possible
 - On segment fault
 - Send UDP packet to master from signal handler
 - Include sequence number of record being processed
 - Skip bad records
 - If master sees two failures for same record, next worker is told to skip the record